


Taiwan Sign Language Translation



Team: 王薌婷 林于湘 林妘鑫 洪睿謙 陳郁婷 張育禎 楊峻棋 賴藍平
Presenter: 洪睿謙 陳郁婷





Demo



久/相見/無
=好久不見





Contents

TSL Translation

01

Introduction

Why we started this project?

02

Data & Methods

How to make this happen?

03

Experiments

What have we done?

04

Conclusion & Prospect

What's next?


01

INTRODUCTION



Motivation

Introduction



Why not...

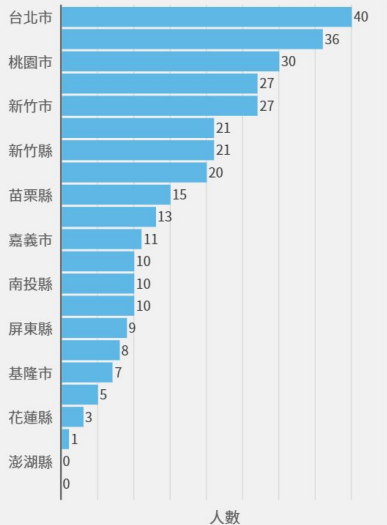
Motivation

Introduction

探究台灣手語翻譯服務現況，思考解決方案。期望改善使用者可能面臨的困境並提升生活品質與便利性。




提供社福服務之手譯員，共324名：所有提供社會福利類手語翻譯類人員，包含乙級、丙級及其他未取得證照但符合規定者。



Goal

Introduction



TSL Translation

Goal

Introduction

- 手語隱含文化獨特性
- 語言系統與時俱進
- Hand gesture recognition
- Hand pose estimation & tracking

新技術+AI = 持續性的研究與學習價值





02



Data & Methods

Data Collection

Data & Methods

手語是多重且連續性的資料，包含：

1. 手形
2. 位置
3. 動作型態
4. 方向
5. 表情



Data Collection

Data & Methods

來源	參考 Youtube 手語教學，自行拍攝
內容	涵蓋六大主題常用語 「年齡、問候、家人、時間、住家、描述」
筆數	共 55 個字詞，可組合至少 25 個短句
特徵	保留手形、位置(胸部以上)、動作型態及方向

有 買 久 一 你 誰 租 棒 零 我 無

今天 昨天 明天 台北 桃園 手語 朋友
什麼 將近 父母 一共 家裡 他們 吃飯
房子 銀行 認識 見她 比較 是嗎 一樣
生日 天氣 上課 孩子 相見 運動 年齡
名字

星期一 星期二 星期三 星期四 星期五
星期六 星期天 捷運站 高鐵到 我問你
還沒有 完了嗎

我們兩個 幾月幾號 會不會呢

「你叫什麼名字」
你 名字 什麼

Data Collection

Data & Methods

我問你			...	
手語			...	
會不會呢			...	

我問你/手語/會不會呢
= 請問你會不會打手語

55 words
60 frames (2 sec * 30 fps)

Hands & Pose Tracking Solution

Data & Methods

Hands


- Palm detection model: [TFLite model](#), [TF.js model](#)
- Hand landmark model: [TFLite model](#), [TFLite model \(sparse\)](#), [TF.js model](#)
- [Model card](#), [Model card \(sparse\)](#)

Pose

- Pose detection model: [TFLite model](#)
- Pose landmark model: [TFLite model \(lite\)](#), [TFLite model \(heavy\)](#)
- [Model card](#)

Holistic

- Hand recrop model: [TFLite model](#)



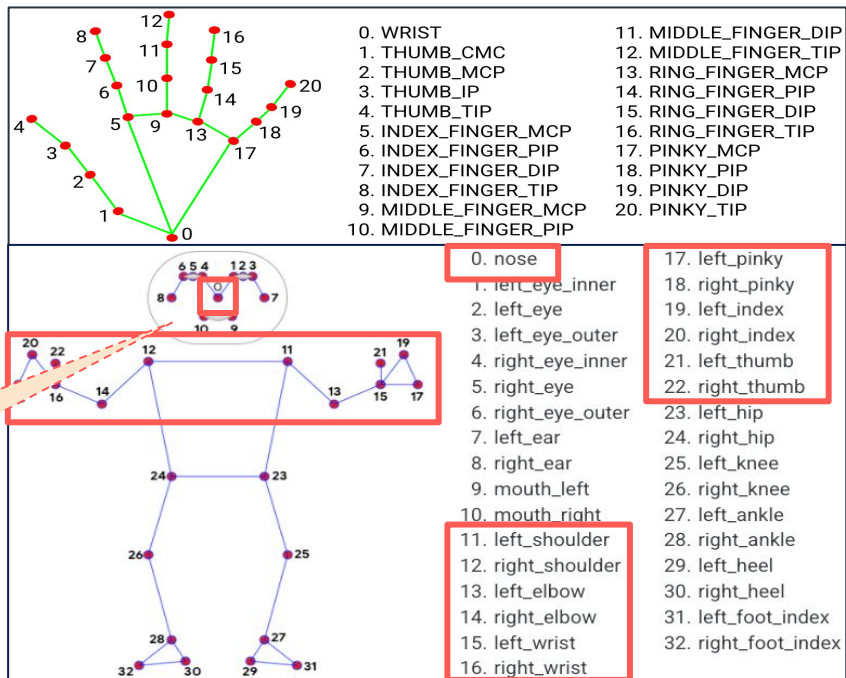
```
Handedness: [classification {  
  index: 0  
  score: 0.99946386  
  label: "Left"  
},  
  classification {  
  index: 1  
  score: 0.99999803  
  label: "Right"  
}]  
hand_landmarks: landmark {  
  x: 0.5342746  
  y: 0.57851875  
  z: -4.123216e-05  
}  
landmark {  
  x: 0.51628554  
  y: 0.5465713
```

Data Preprocessing

- Remove unused landmarks {face, pose 1~10, 23~32} and visibility
- Keep **55 landmarks** {right & left hands, pose 0, 11~22}
- Each has xyz axis
- $55 * 3 = 165$ features

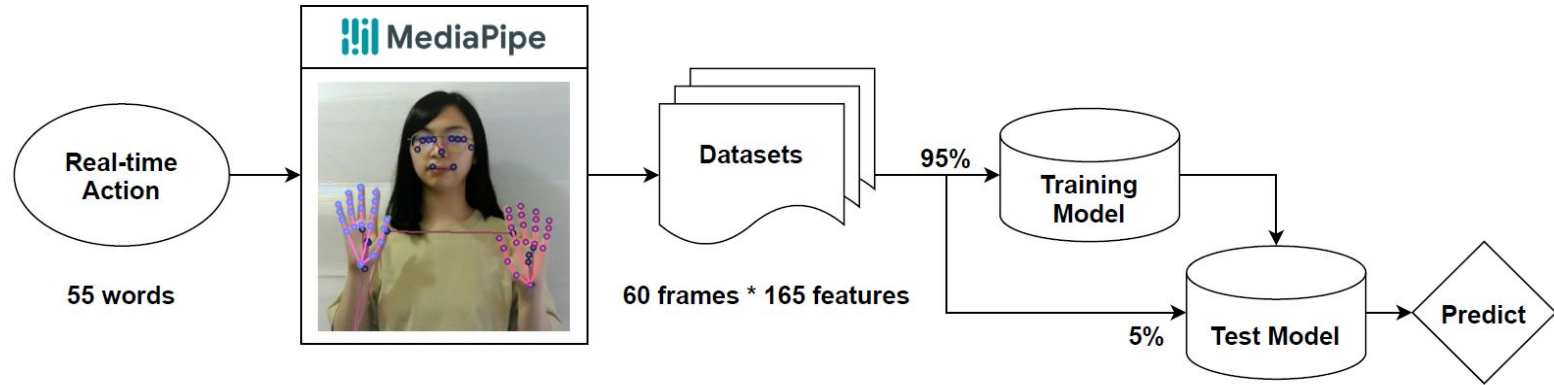
以鼻子為中心
轉換相對座標

Data & Methods



Workflow

Data & Methods




Filming -----> **Array** ---> **AI Model** --> **Words**

5 ppl

03

EXPERIMENTS



Our Journey Begins Here...

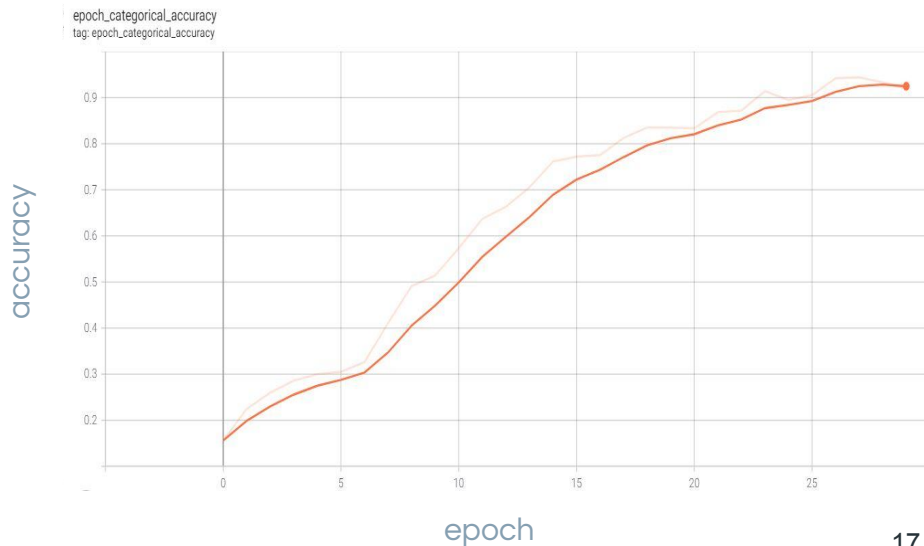
1. Small datasets to test models (10 words)
2. Began with famous RNN model - **LSTM**
3. Turned to **GRU** model to increase speed



	LSTM	GRU
Train accuracy	0.7333	0.9193
Train loss	0.7887	0.2065
Test accuracy	0.7333	0.9000
Test loss	0.8419	0.1724


Experiments

GRU



Sticking with GRU...

4. Bigger datasets based on GRU model



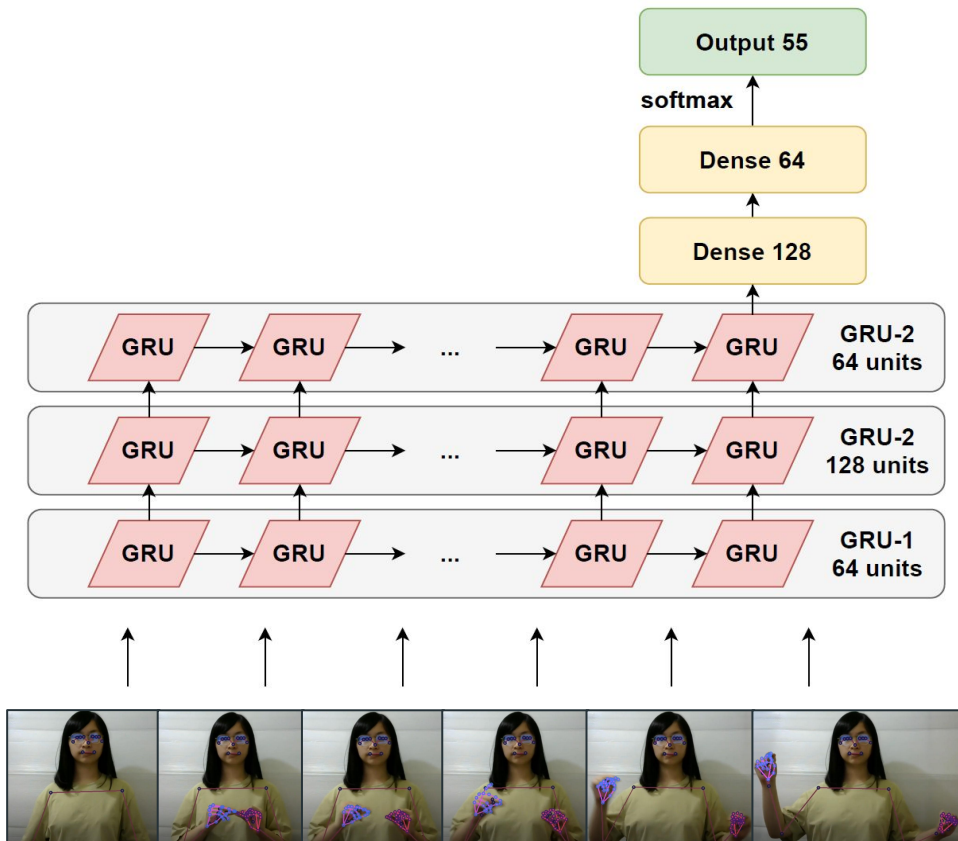
	GRU-10	GRU-55
Epoch	30	75
Train accuracy	0.9193	0.9573
Train loss	0.2065	0.1321
Test accuracy	0.9000	0.8909
Test loss	0.1724	0.4429
Real-time test	60% (6/10)	42% (23/55)

Experiments

More
experiments
based on
GRU model

GRU Model Architecture

Experiments



Input shape (60, 165)



Boss 1

GRU model no longer performs
as well with 55 words

What have we done?

5. Chose **ELU** instead of ReLU to preserve features



	ReLU	ELU
Train accuracy	0.9420	0.9910
Train loss	0.1960	0.0420
Test accuracy	0.8850	0.9330
Test loss	0.3750	0.2030
Real-time test	55% (30/55)	67% (37/55)

Experiments

Model: "sequential_1"

Layer (type)	Output Shape	Param #
gru_3 (GRU)	(None, 60, 128)	113280
gru_4 (GRU)	(None, 60, 256)	296448
gru_5 (GRU)	(None, 128)	148224
dense_3 (Dense)	(None, 256)	33024
dense_4 (Dense)	(None, 128)	32896
dense_5 (Dense)	(None, 55)	7095
Total params: 630,967		
Trainable params: 630,967		
Non-trainable params: 0		

Experiments

***Discovered a BUG in Coordinate
Transformation***

BREAKING

NEWS



Boss 2

Hand coordinates was not transformed into relative coordinate !!


What's wrong?




	ELU Old data	ELU New data?
Train accuracy	0.9910	0.9911
Train loss	0.0420	0.0389
Test accuracy	0.9330	0.9333
Test loss	0.2030	0.2127
Real-time test	67% (37/55)	51% (28/55)

Experiments

We looked
into the data

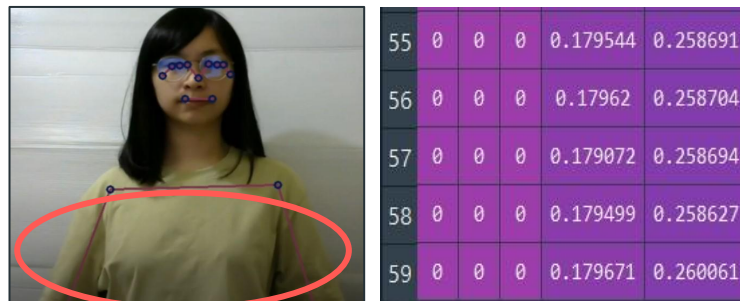


What's wrong?

	ELU Old data	 ELU New data?
Train accuracy	0.9910	0.9911
Train loss	0.0420	0.0389
Test accuracy	0.9330	0.9333
Test loss	0.2030	0.2127
Real-time test	67% (37/55)	51% (28/55)


Experiments

When hand isn't detected
it'll be filled with 0




What have we done?

6. Implemented **mask** to preserve "0"



Experiments

Only keep the results
where the value is true
(0 - nose = obtain value)




Tuning Model

7. Increase model size (add to GRU 256,512)



	New data (128, 256)	New data (256, 512)
Train accuracy	0.9810	0.9880
Train loss	0.0640	0.0410
Test accuracy	0.9390	0.9450
Test loss	0.2110	0.1700
Real-time test	67% (37/55)	75% (41/55)

Experiments





Boss 3

Ceiling: reach the
maximum potential


What have we done?

8. Increase model size (add to GRU 512,1028)
9. Increase number of videos (add to 100)



	v=60 (512, 1028)	v=100 (512, 1028)
Train accuracy	0.9940	0.9970
Train loss	0.0320	0.1890
Test accuracy	0.9650	0.9820
Test loss	0.1090	0.0530
Real-time test	76% (42/55)	87% (48/55)

Experiments



Hyperparameters Fine Tuning

Experiments

10. Used **AutoKeras** instead of grid search

RNNBlock

```
autokeras.RNNBlock(return_sequences=False,
bidirectional=None, num_layers=None,
layer_type=None)
```


DenseBlock

```
autokeras.DenseBlock(num_layers=None,
num_units=None, use_batchnorm=None,
dropout=None)
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 60, 165)]	0
cast_to_float32 (CastToFloat)	(None, 60, 165)	0
bidirectional (Bidirectional)	(None, 60, 330)	328680
bidirectional_1 (Bidirectional)	(None, 330)	492030
dense (Dense)	(None, 64)	21184
batch_normalization (Batch Normalization)	(None, 64)	256
re_lu (ReLU)	(None, 64)	0
dense_1 (Dense)	(None, 55)	3575
classification_head_3 (Softmax)	(None, 55)	0
Total params: 845,725		
Trainable params: 845,597		
Non-trainable params: 128		

Hyperparameters Fine Tuning

10. Used **AutoKeras** instead of grid search



	v=100 (512, 1028)	AutoKeras
Train accuracy	0.9970	0.9991
Train loss	0.189	0.0270
Test accuracy	0.982	0.9927
Test loss	0.053	0.0504
Real-time test	87% (48/55)	95% (52/55)

Experiments

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 60, 165)]	0
cast_to_float32 (CastToFloat)	(None, 60, 165)	0
bidirectional (Bidirectional)	(None, 60, 330)	328680
bidirectional_1 (Bidirection	(None, 330)	492030
dense (Dense)	(None, 64)	21184
batch_normalization (BatchNo	(None, 64)	256
re_lu (ReLU)	(None, 64)	0
dense_1 (Dense)	(None, 55)	3575
classification_head_3 (Softm	(None, 55)	0
Total params: 845,725		
Trainable params: 845,597		
Non-trainable params: 128		

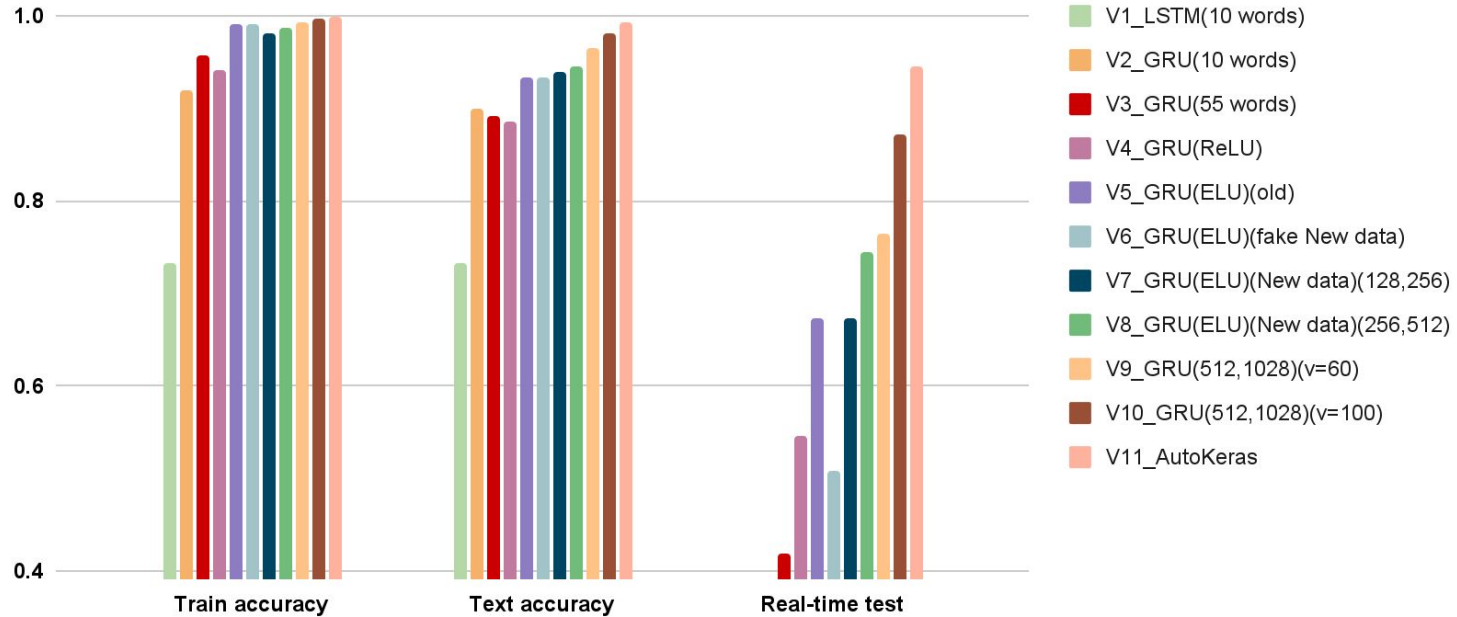


04

CONCLUSION & PROSPECT


Review: Comparison of All Model

Conclusion & Prospect



Review: Comparison of Datasets

Conclusion & Prospect



More words

More videos

Conclusion



Conclusion & Prospect

- AutoKeras outperforms V10 (GRU 512) on all numerical aspects
- Currently V10 has the upper hand in terms of real-time due to technical difficulties (time)

Conclusion

Conclusion & Prospect

- GRU performs better on smaller datasets



Comparable with previous TSL project

	Previous	Ours
Words	7	55 ↑
Number of videos	200	100 ↓
Frames	90	60 ↓
Model	LSTM	GRU
Best test accuracy	0.9963	0.9927


Conclusion & Prospect

Real-time test:

95%
(52/55)

Confusion Matrix

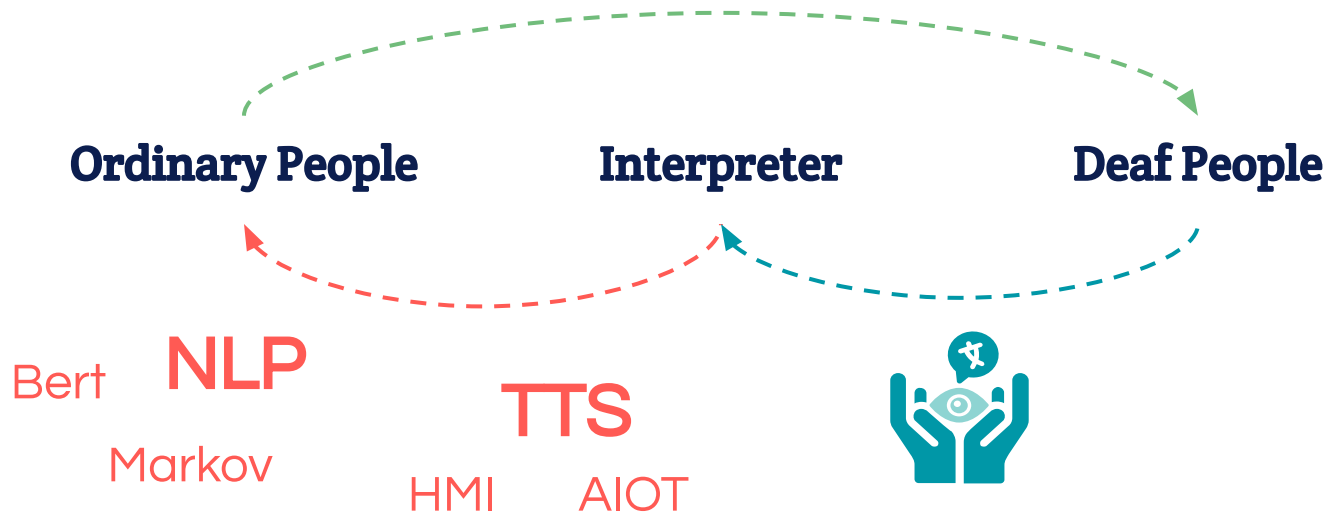
Conclusion & Prospect



Prospect

Conclusion & Prospect

Education



+ Domain Knowledge



Thanks!

Taiwan Sign Language Translation

王薌婷 林于湘 林妘鑫 洪睿謙 陳郁婷 張育禎 楊峻棋 賴藍平

CREDITS: This presentation template was created by Slidesgo, Flaticon, Freepik and Storyset